

Intro to RL in Julia


RL == Reinforcement Learning (not Rugby League)

In the beginning...



Programmers




















Reinforcement Learning



Key Features of RL

- Agent interacting with an environment (Trial and Error)
 - sequence of states, actions, rewards

Key Features of RL

- Agent interacting with an environment (Trial and Error)
 - sequence of states, actions, rewards
- Feedback not necessarily immediate (Delayed Rewards)

Key Features of RL

- Agent interacting with an environment (Trial and Error)
 - sequence of states, actions, rewards
- Feedback not necessarily immediate (Delayed Rewards)
- Vs Supervised ML:
 - single decision
 - immediate/“dense” feedback

Key Features of RL

- Agent interacting with an environment (Trial and Error)
 - sequence of states, actions, rewards
- Feedback not necessarily immediate (Delayed Rewards)
- Vs Supervised ML:
 - single decision
 - immediate/“dense” feedback
- Vs Unsupervised ML: there is *some* feedback

Examples where RL is used

- **Robotics**

- Walking
- Picking things up, Pouring a coffee, Making a sandwich
- State: sensory input
- Action: torques to apply to motors

- **A/B Testing**

- State: URL, search history, age, interests, etc.
- Action: which ad to show, where to place it

Where RL is used

- **Data Center Cooling*:**
 - State: temperatures, power, pump speeds, etc.
 - Action: where to pump coolant
- **Taxi/Ride Share Dispatch**
 - State: locations of taxis and riders
 - Action: dispatch taxi X to location Y
- **Game Playing AI:** AlphaGo/Zero, Backgammon, Atari Games, DOTA2, StarCraft2

Key Idea: Value & Credit Assignment

- Rewards are often delayed, so reward in next timestep \neq value of action
 - I.e. hard to know which action(s) to credit for a particular reward
 - E.g. hitting a good shot in tennis - don't score immediately
 - Extreme case: decisive move in chess, reward only received at end of game
- Key component of RL methods is estimating the **Value** of an action
- Defined as
 - $V(s)$: Sum of (expected) future rewards starting in state s
 - $Q(s,a)$: Sum of (expected) future rewards starting in state s , taking action a
 - (and then continuing with current policy)

Model Based vs Model Free

- Model Free
 - Learn value of states, or state, action pairs
- Model Based
 - explicit environment model
 - model errors accumulate
 - less explored, much potential

Tabular RL

Table storing, for every state, action pair (e.g. all possible player/dealer hand combinations in blackjack):

- Action to take in that state (hit, stay, double, split)
- (Expected) Value of each state: (prob of winning * amount that will be won) - stake


Value Iteration: run a lot of simulations, being sure to try all actions in each state - will give you a “monte carlo” estimate of value.



Tabular RL

There are other algorithms that are more efficient (than Monte-Carlo Value Iteration), but they're not generally scalable to large state/action spaces.

Deep RL



Limitations and Research Frontiers

- Simulators required
 - sample inefficient
 - unsafe
- Exploration (vs Exploitation)
 - Intrinsic reward - e.g. curiosity based
- Model Based
- Low level control
- Multi-task learning
- ...

Why Julia

- Speed: <https://github.com/JuliaML/OpenAIGym.jl/pull/13>
 - FFI: use C, Python libs easily and performantly
 - C++ not as straightforward but do-able and improving
- Same language for research and deployment
- Distributed/Concurrent/Parallel execution in the stdlib
 - Great GPU support: CuArrays, CUDAnative
- Package ecosystem:
 - JuliaDiffEq.jl
 - Optim.jl and other Optim packages (NLOpt.jl, BlackBoxOptim.jl, ...)
 - Distributions.jl - Probability distributions
 - Enabled by: multiple dispatch and AbstractArray

Julia Reinforcement Learning

(Code) <https://github.com/JuliaReinforcementLearning/ReinforcementLearning.jl>

(Docs) <https://juliareinforcementlearning.github.io/ReinforcementLearning.jl/latest/>

RL in Julia Roadmap

- ReinforcementLearning.jl v0.3.0 - see <https://github.com/JuliaReinforcementLearning/ReinforcementLearning.jl/issues/24>
- Distributed Execution
- More environments
 - Integrate with JuliaRobotics Stack
 - Faster simulations
 - Differentiable environments with Zygote
 - Wrap more environments with C/C++

Resources

- Julia RLish stuff:
 - <https://github.com/JuliaReinforcementLearning/ReinforcementLearning.jl>
 - Julia Robotics talk @ JuliaCon 2018 (not RL, but related + we hope to integrate as RL envs) <https://www.youtube.com/watch?v=dmWQtI3DFFo>
- Recommended RL reading/resources:
 - <http://karpathy.github.io/2016/05/31/rl/>
 - David Silver (DeepMind) Deep RL lectures:
<https://www.youtube.com/playlist?list=PLqYmG7hTraZDM-OYHWgPebj2MfCFzFObQ>
 - OpenAI spinning up in Deep RL: <https://spinningup.openai.com> (project idea: translate this into Julia)
 - *The RL Textbook*: <http://incompleteideas.net/book/the-book-2nd.html>